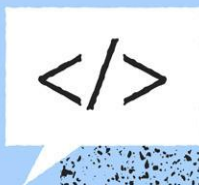




Никита Гришко / Lead Software Engineer

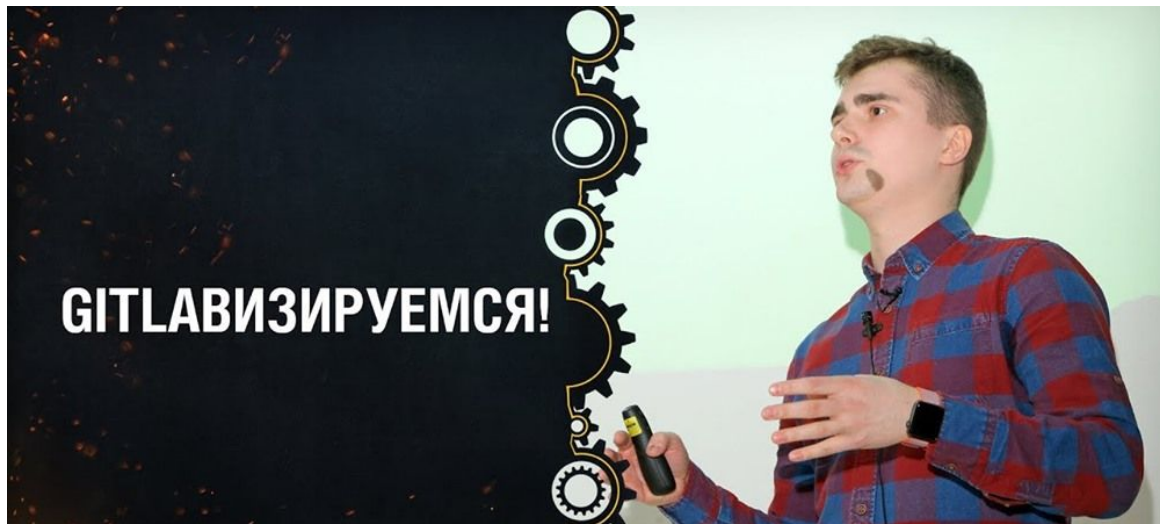
Kubernetes- зируемся



whoami

- Lead Software Engineer at PandaDoc
- 10 years experience in backend development
- Moved from product development to platform engineering

A few years ago...



<https://www.youtube.com/watch?v=QhdggHKnVtk>

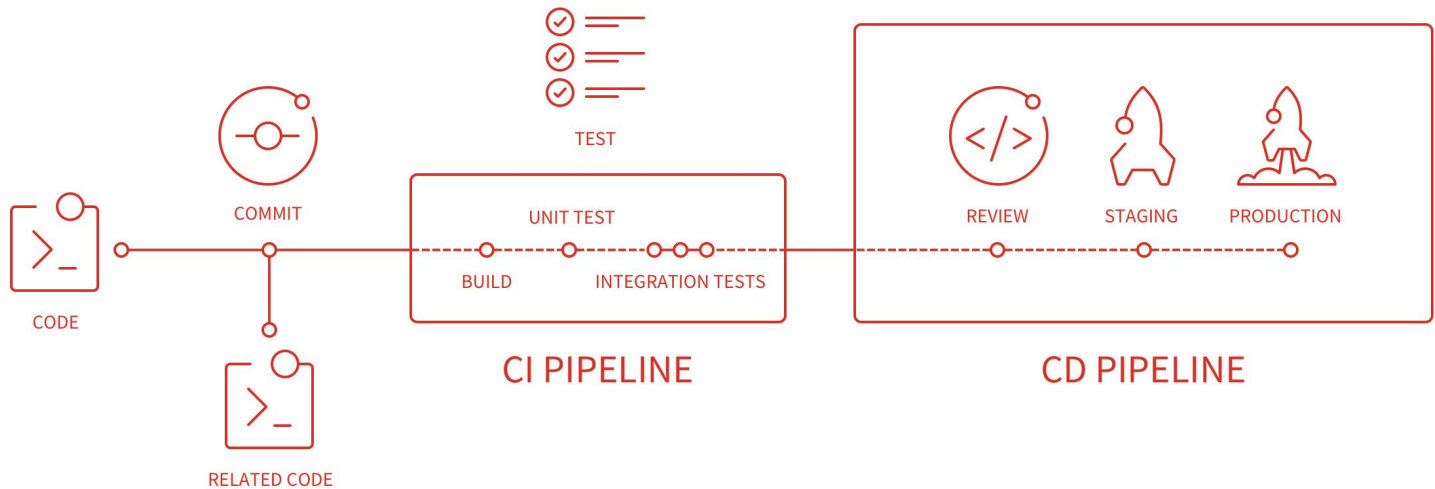
A few years ago...

- AWS EC2 + Docker
- Gitlab for source code management and CI
- First steps in continuous integration

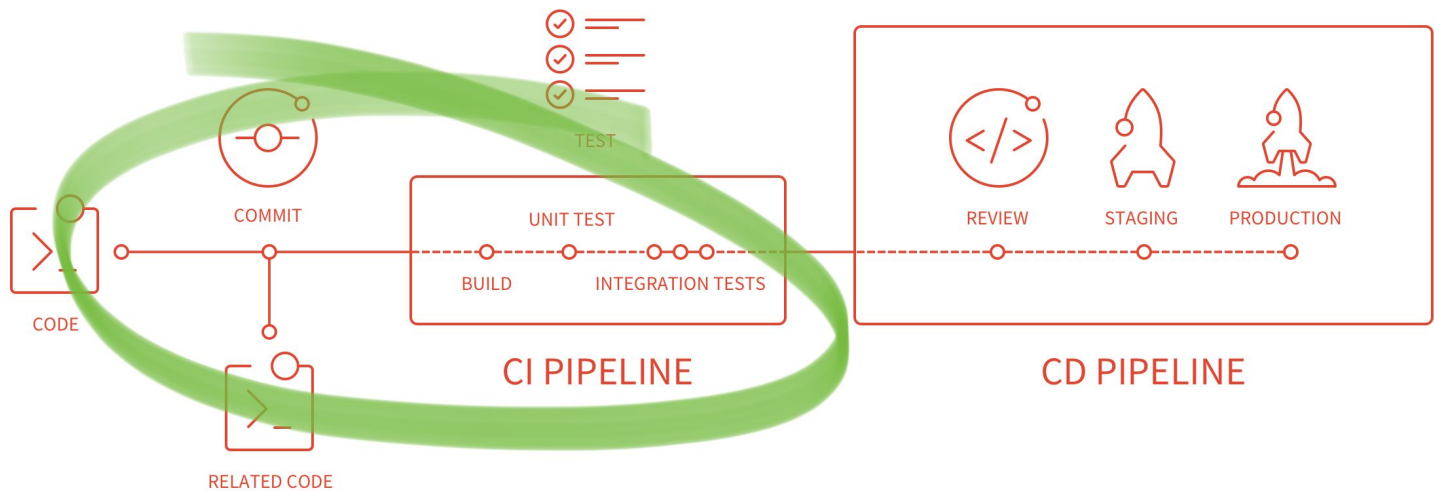


<https://www.youtube.com/watch?v=QhdggHKnVtk>

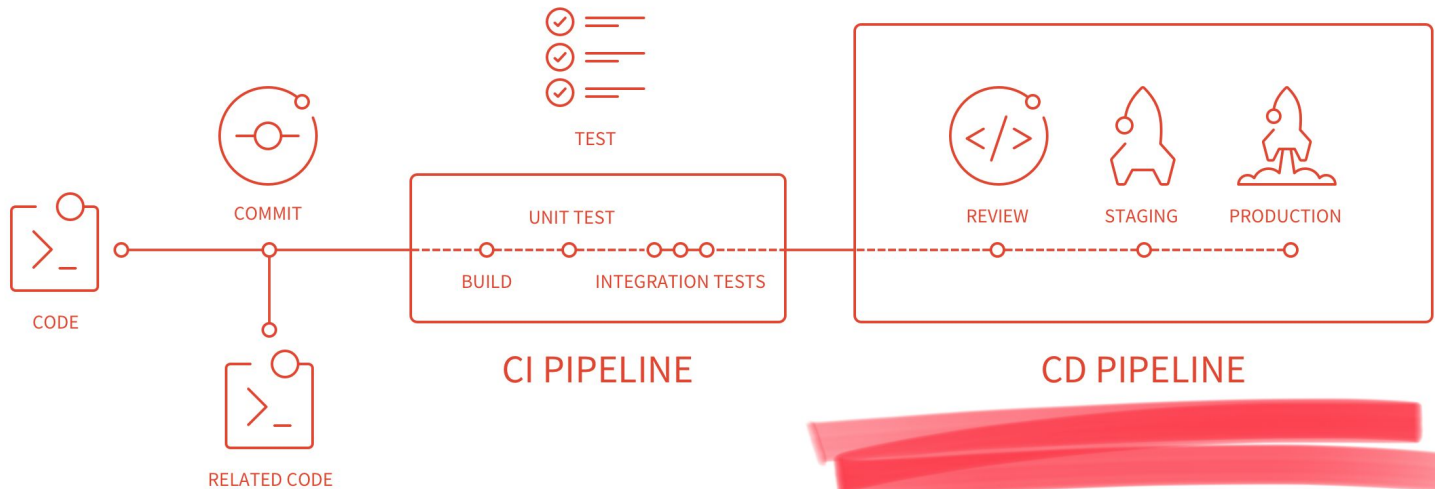
A few years ago...



A few years ago...



A few years ago...




No, no, no, no

- No bu****it bingo
- No manuals retelling
- No “silver bullet” solutions
- No complex technical solutions

Yes!

- Kubernetes? Why?
- From AWS EC2 to AWS EKS: our step by step guide
- Our CI/CD pipeline

Still not interested yet?

- More than 50 services migrated in 5 months (including legacy ones)
- Over 20+ environments run on Kubernetes 
- No epic fails, please believe me

Still not interested yet?

- More than 50 services migrated in 5 months (including **legacy** ones)

Using Kubernetes for Legacy services.



Kubernetes? Why?



Artem Anokhin

@SumLare



Инвесторы таки дали денег.

Что видят менеджеры: фух, ещё пару кварталов протянем

Что видят разработчики: пора пробовать k8s с кафкой, пока не поздно

Kubernetes? Why?

- Terraform + Ansible



Kubernetes? Why?

- Terraform + Ansible:
 - ~~Too difficult and no one wants it~~
 - A lot of repositories to put changes
 - Hard and painful rollback
 - Anyone from another team can accidentally break your deployment configuration
 - 2-4 days for initial service setup
 - **Bad cycle time experience**

Kubernetes? Why?

- Terraform is slow...
 - Unexpected load on your service? Want to increase nodes count or change node type?
 - **Again bad cycle time experience**

Kubernetes? Why?

- Ansible is slow too...
 - Install system packages
 - Setup users
 - Render templates
 - Pull and run containers
 - **And again bad cycle time experience**

Kubernetes? Why?

- Ansible is slow too...



This run spent:

- 2 ms waiting;
- 1 min 18 sec build duration;
- 1 min 18 sec total from scheduled to completion.



This run spent:

- 4 ms waiting;
- 6 min 19 sec build duration;
- 6 min 19 sec total from scheduled to completion.

VS

deploy-production

[Retry](#)

Duration: 1 minute 45 seconds

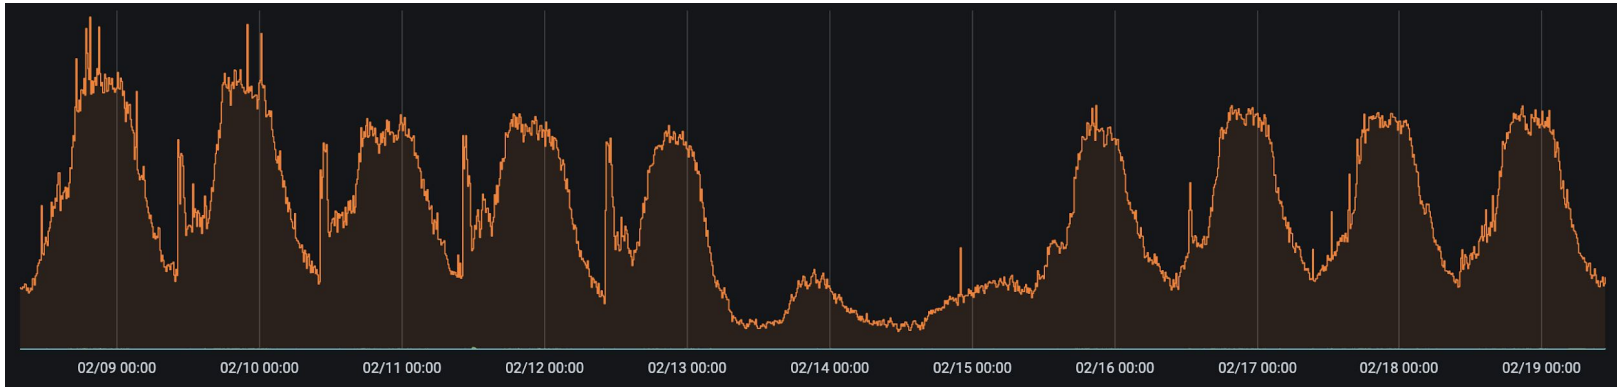
Timeout: 1h (from project)



Runner: gitlab-runner-prod-main-shared-55b7c64b59-257dm (#612)

Tags: [shared](#) [prod-main](#)

Kubernetes? Why?



Time-dependent load on services

Kubernetes? Why?

- Time-dependent load on services
- Bad resources utilization
- Risk of downtime in case of unexpected load

So why? What we want?

- Everything in one repository: source code, deployment configuration, even dashboards and alerts
- (Auto)scaling and resources optimization
- Cost optimization
- **Better cycle time and development experience**

Maybe XYZ?

- Kubernetes de facto industry standard
- A lot of experience and tools in community



Ready for maintenance standard

Ready for maintenance standard?

- Only Docker images
- Scripts: migrate.sh and entrypoint.sh
- Configuration using environment variables
- Logs to STDOUT in JSON
- Secrets in HashiCorp Vault
- Expose 4284 port for health check
- Prometheus and standardized set of metrics
- Transports: NATS and Kafka
- ...

Ready for maintenance standard?

- gw* (gwjava-*, gwpy-*)
- ms* (msjava-*, mspy-*, msscala-*)

Ready for maintenance standard!

Standard allows us to unify services and simplify deployments



Helm charts

- Designed by aliens for predators...

```

{{- $deploymentFullname := printf "%s-%s" (include "service.fullname" .) .Release.Name }}
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ $deploymentFullname }}
  labels:
    {{- include "service.labels" $ | nindent 4 }}
spec:
  replicas: {{ $deploymentOptions.replicas }}
  strategy:
    {{- with $deploymentOptions.strategy }}
    {{- toYaml . | nindent 4 }}
    {{- else }}
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
    type: RollingUpdate
    {{- end }}
  selector:
    matchLabels:
      {{- include "service.selectorLabels" $ | nindent 6 }}
  template:
    metadata:
      labels:
        {{- include "service.selectorLabels" $ | nindent 8 }}
        {{- include "service.deploymentLabels" $ | nindent 8 }}
      app-deployment: {{ $deploymentFullname }}
      app-deployed-at: {{ now | unixEpoch | quote }}
    spec:
      {{- with $deploymentOptions.affinity }}
      affinity:
        {{- toYaml . | nindent 8 }}


```

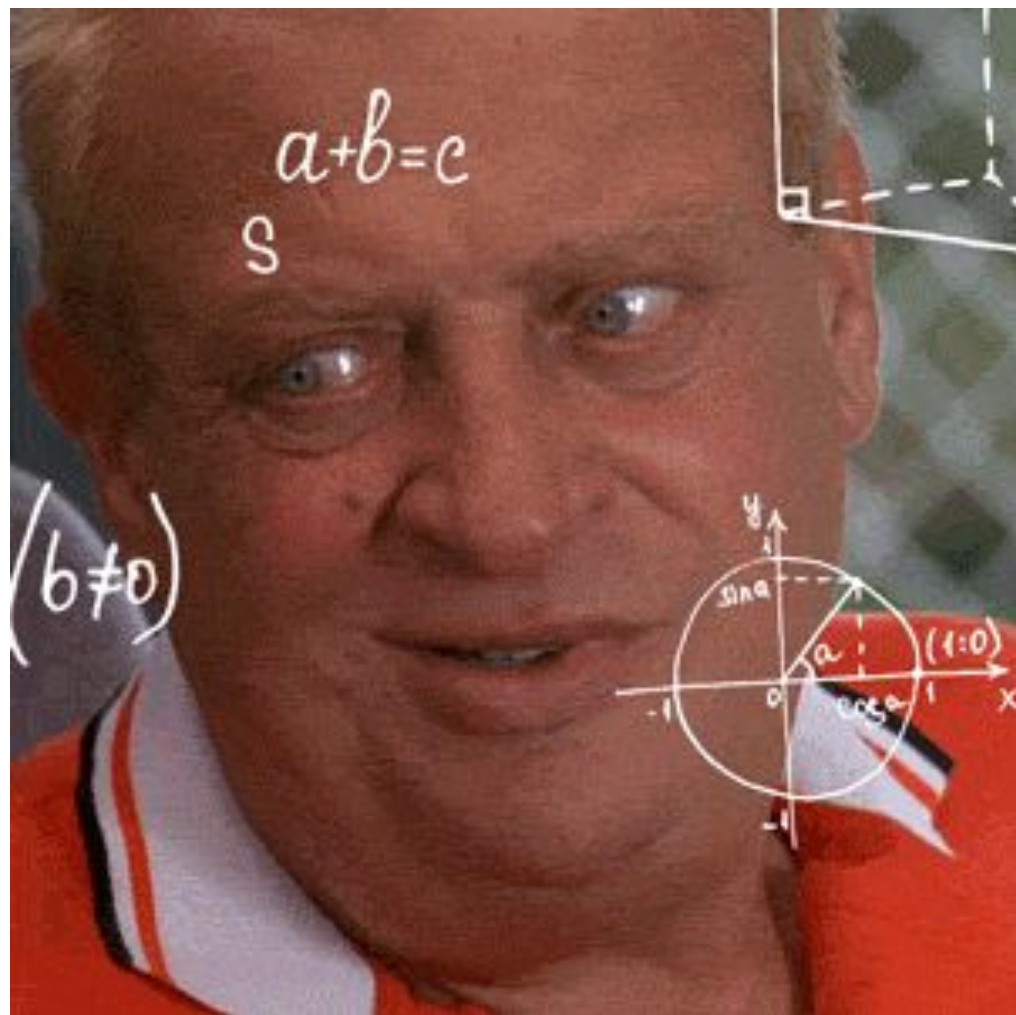
```

{{/*
Create a default fully qualified app name.
We truncate at 63 chars because some Kubernetes name fields are limited to this (by the DNS naming spec).
If release name contains chart name it will be used as a full name.
*/}}
{{- define "service.fullname" -}}
{{- if .Values.fullnameOverride }}
{{- .Values.fullnameOverride | trunc 63 | trimSuffix "-" }}
{{- else }}
{{- $name := default .Chart.Name .Values.nameOverride }}
{{- if contains $name .Release.Name }}
{{- .Release.Name | trunc 63 | trimSuffix "-" }}
{{- else }}
{{- printf "%s-%s" .Release.Name $name | trunc 63 | trimSuffix "-" }}
{{- end }}
{{- end }}
{{- end }}

```

Helm charts

- Designed by aliens for predators...
- YAML + Golang template engine 
- Too complex and too easy to make a mistake
- **Can we provide a better experience?**





Helm charts

- Shared Helm chart for the whole company
- Expose Helm values only for engineers
helm/values/{defaults|preprod|prod}.yaml

```
1  deployments:
2    service:
3      replicas: 1
4      containers:
5        example:
6          image:
7            repository: product/mspy-example
8            command:
9              - ./entrypoint.sh
10
11  secrets:
12    env-secrets:
13      data:
14        APP_NATS:
15          servers:
16            - nats://user:password@host:port
17
```

PandaDoc chart

- Readiness/liveness probes
- CPU/Memory resources
- Ports
- Network services
- Service monitors
- ...



PandaDoc chart

Thanks to **ready for maintenance standard** we know everything about our services and we can provide **useful defaults**

```
1 migrations:
2   container:
3     image:
4       repository: product/mspy-example
5     command:
6       - ./migrate.sh
7
```

```
1 migrations:
2   container:
3     image:
4     repository: product/mspy
5     command:
6     - ./migrate.sh
7
```

```
1 cronJobs:
2   job-cleanup:
3     container:
4       image:
5       repository: product/mspy-example
6       command:
7       - ./cleanup.sh
8     schedule: "@monthly"
9
```

```
1 migrations:
2   container:
3     image:
4       repository: product/mspy
5     command:
6       - ./migrate.sh
7
```

```
1 cronJobs:
2   job-cleanup:
3     container:
4       image:
5         repository: product/mspy-example
6       command:
7         - ./cleanup.sh
8     schedule: "@monthly"
```

```
1 jobs:
2   job-cleanup:
3     container:
4       image:
5         repository: product/mspy-example
6       command:
7         - ./cleanup.sh
8
```



```
1 migrations:
2   container:
3     image:
4     repository: product/mspy
5     command:
6     - ./migrate.sh
7
```

```
1 cronJobs:
2   job-cleanup:
3     container:
4     image:
5     repository: product/mspy-example
6     command:
7     - ./cleanup.sh
8   schedule: "@monthly"
```

```
1 jobs:
2   job-cleanup:
3     container:
4     image:
5     repository: product/mspy-e
6     command:
7     - ./cleanup.sh
8
```

```
1 ingresses:
2   external:
3     host: random.pandadoc.com
4     annotations:
5     nginx.org/websocket-services: mspy-example
6     paths:
7     - deployment: service
8     path: /
9     port: 80
10
```





Gitlab CI/CD

- Build Docker image
- Run Helm (upgrade, history, ...)
- Verify that all pods are ready
- Send notifications (Slack, ...)
- Annotate Grafana dashboards
- ...



Gitlab CI/CD

- Ctrl-C / Ctrl-V is not an option
- How can we control versions of tools we use?



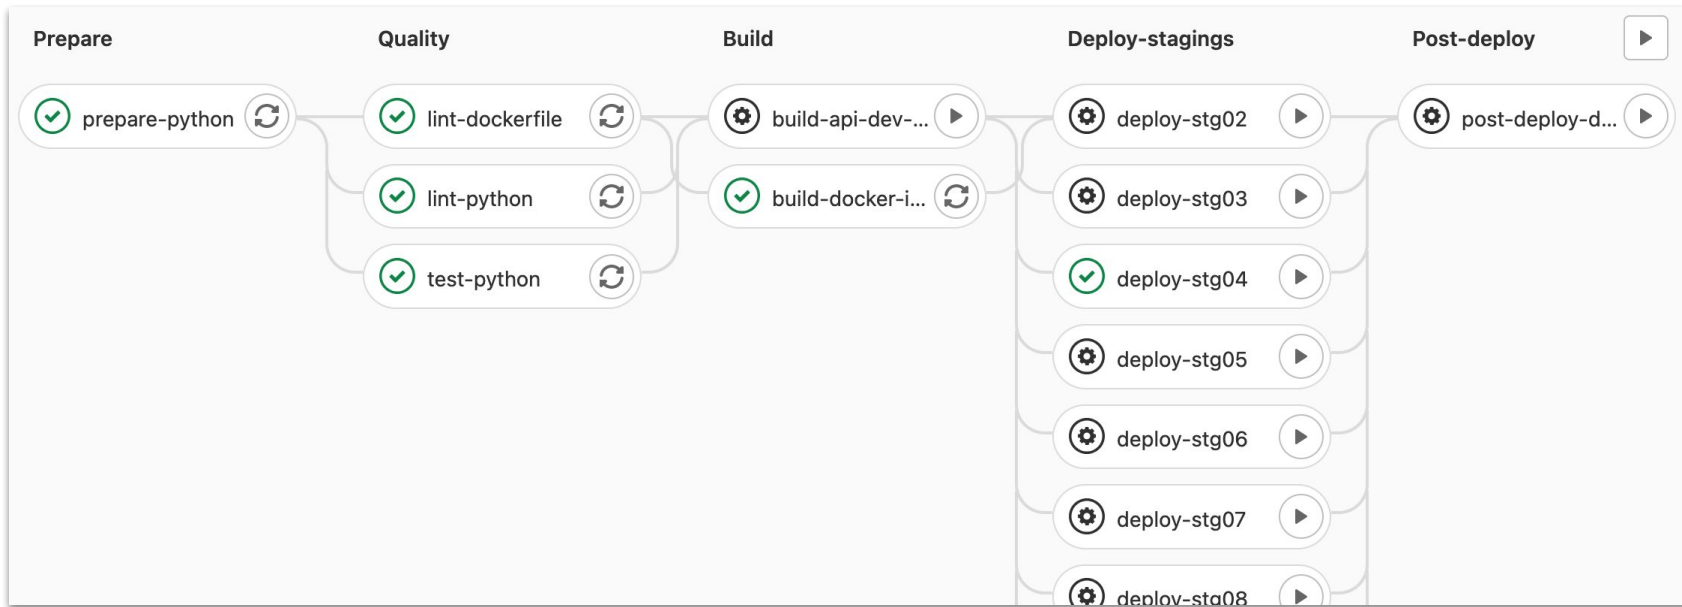
Gitlab CI/CD

- <https://docs.gitlab.com/ee/ci/yaml/README.html#include>
- You can use *include* to include external YAML files in your CI/CD configuration

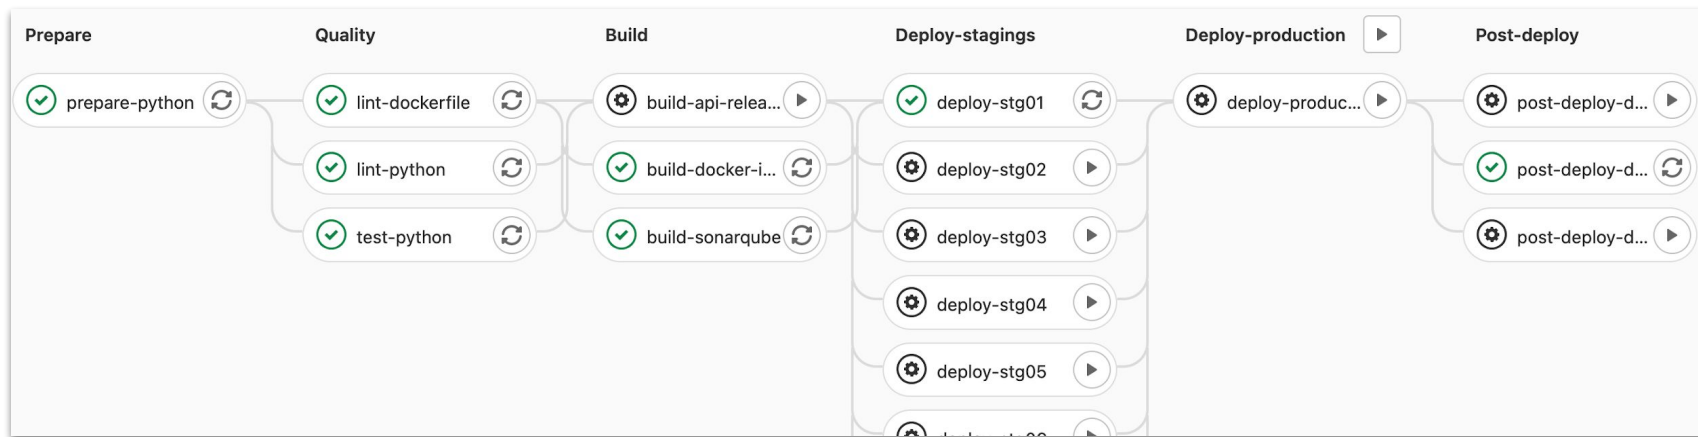


```
1  include:
2    - project: platform/ci/k8s-pandakube
3      ref: master
4      file: pipelines/deploy-migrate.yml
5    - project: platform/ci/k8s-pandakube
6      ref: master
7      file: pipelines/post-deploy-dashboards.yml
8    - project: platform/ci/pipelines
9      ref: master
10     file: docker/uber.yml
11    - project: platform/templates/template-pdms-service
12      ref: master
13      file: pipelines/uber.yml
14
15  stages:
16    - prepare
17    - quality
18    - build
19    - deploy-stagings
20    - deploy-production
21    - post-deploy
22
23  variables:
24    CI_DEPLOY_GRAFANA_PROVISION_MS_APP_METRICS_DASHBOARD: "true"
25    CI_PDMS_SERVICE_PYTHON_IMAGE_VERSION: "3.8"
26    DOCKER_SERVICE_PIPELINE_IMAGE: $ARTIFACTORY_DOCKER_REGISTRY/$CI_PROJECT_PATH:$CI_COMMIT_REF_SLUG-$CI_PIPELINE_ID
27
```













Gitlab CI/CD (feature branch)



Gitlab CI/CD (master branch)



Gitlab CI/CD (rollbacks)

prod							Monitoring	Edit	Stop
Status	ID	Triggerer	Commit	Job	Created	Deployed			
success	#2263		 master  d9f8f188  PLT-90: [tech] Update resources f...	deploy-produ...	1 day ago	1 day ago	 		
success	#2242		 master  d9f8f188  PLT-90: [tech] Update resources f...	deploy-produ...	6 days ago	6 days ago	 	<div>Rollback environment</div>	



Gitlab CI/CD

- k8s-pandakube — our golden image for Kubernetes deployments
- Pinned versions of Helm and kubectl
- Helm charts
- Grafana dashboards
- ...

(Auto)scaling

(Manual)scaling

Manual scaling is easy as possible and can be applied in a few minutes (depends on the time of pipeline)

```
1  deployments:
2    service:
3      replicas: 10
4      containers:
5        commentator:
6          resources:
7            limits:
8              cpu: 500m
9              memory: 512Mi
10           requests:
11             cpu: 500m
12             memory: 512Mi
13
```

(Auto)scaling

- Horizontal Pod Autoscaler
- <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>
- CPU / Memory



(Auto)scaling

- Horizontal Pod Autoscaler
- <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>
- CPU / Memory
- **We want to scale by Prometheus metrics**



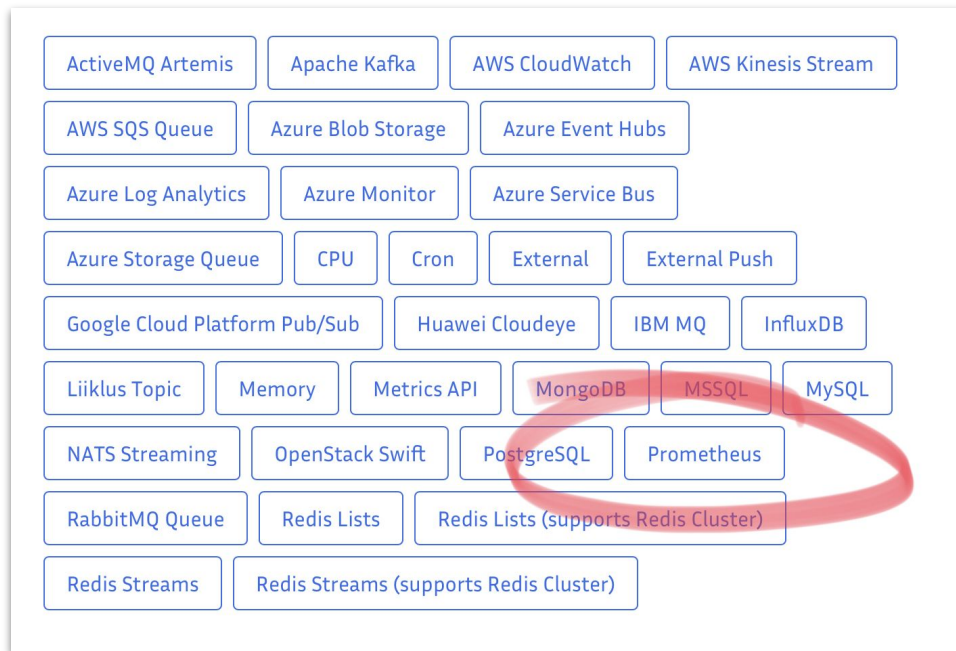
(Auto)scaling

- KEDA — Kubernetes Event-driven Autoscaling
- <https://keda.sh/>
- KEDA works alongside standard Kubernetes components like the Horizontal Pod Autoscaler and can extend functionality without overwriting or duplication



(Auto)scaling

Available scalers for KEDA 2.2



(Auto)scaling

```
1  deployments:
2    service:
3      replicas: 1
4      autoscaling:
5        maxReplicas: 10
6        triggers:
7          - source: prometheus
8            metricName: MetricName
9            query: sum(rate(metric{label="value"}[1m]))
10           threshold: 0.5
11          - source: cpu
12            type: Utilization
13            value: 90
14
```



+



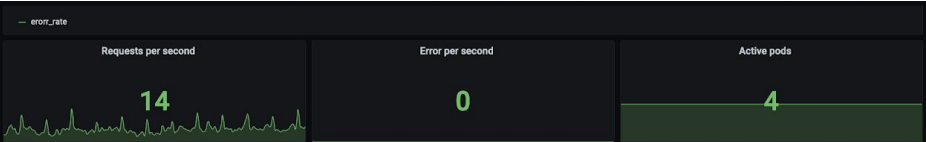
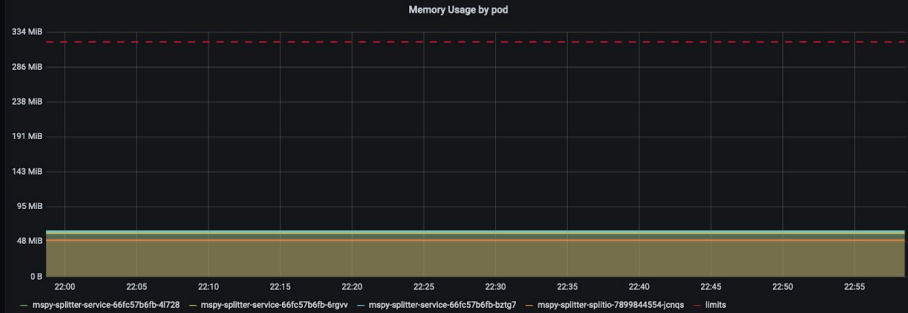
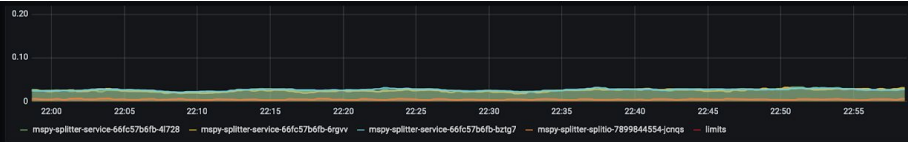
Grafana & Prometheus

- Everything in one repository
- alerts/l2.yaml
- dashboards/queues.json
- Any alerts (as alert rules) or dashboards (as config maps) definitions can be deployed to Kubernetes and applied by **Prometheus operator**

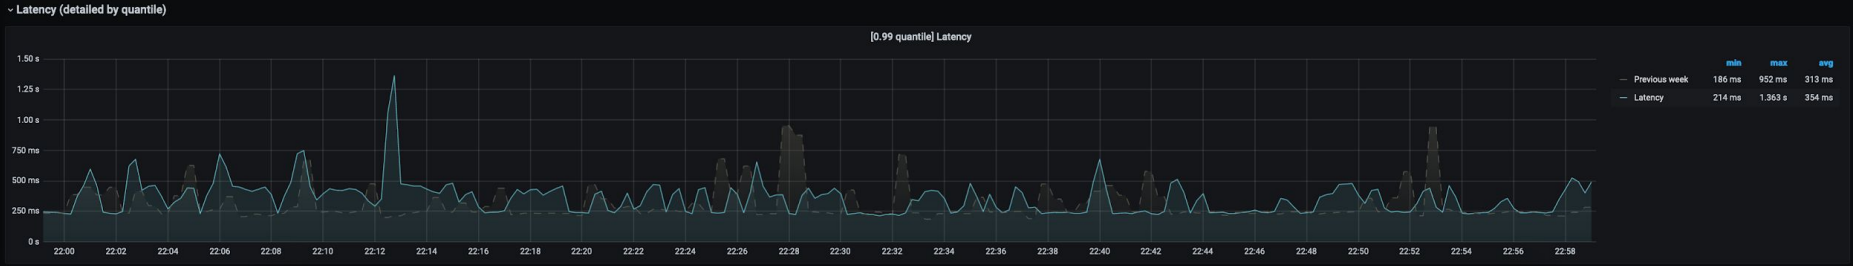
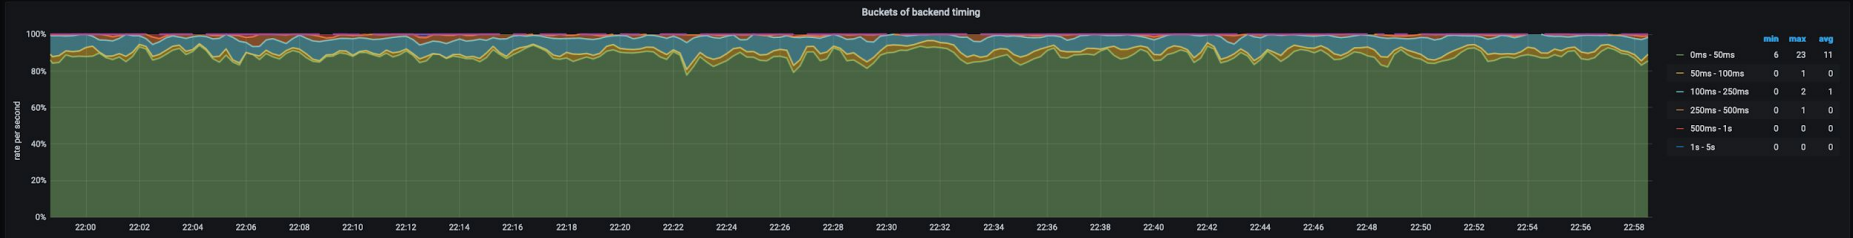
```
1  include:
2    - project: platform/ci/k8s-pandakube
3      ref: master
4      file: pipelines/deploy-migrate.yml
5    - project: platform/ci/k8s-pandakube
6      ref: master
7      file: pipelines/post-deploy-dashboards.yml
8    - project: platform/ci/pipelines
9      ref: master
10     file: docker/uber.yml
11    - project: platform/templates/template-pdms-service
12      ref: master
13      file: pipelines/uber.yml
14
15  stages:
16    - prepare
17    - quality
18    - build
19    - deploy-stagings
20    - deploy-production
21    - post-deploy
22
23  variables:
24    CI_DEPLOY_GRAFANA_PROVISION_MS_APP_METRICS_DASHBOARD: "true"
25    CI_PDMS_SERVICE_PYTHON_IMAGE_VERSION: "3.8"
26    DOCKER_SERVICE_PIPELINE_IMAGE: $ARTIFACTORY_DOCKER_REGISTRY/$CI_PROJECT_PATH:$CI_COMMIT_REF_SLUG-$CI_PIPELINE_ID
27
```

Grafana & Prometheus

- Shared Grafana dashboard: CPU/memory usage; latency by percentiles; RPS; top slowest requests and more



Top 10 slow handlers (by endpoint)			Top 10 error per endpoint		
endpoint	code	Response Time +	endpoint	code	Value
get_treatments	ok	220 ms	get_organization_treatment	error_internal	0 s
get_treatments	ok	204 ms	get_organization_treatment	error_internal	0 s
get_treatment	ok	87.0 ms	get_treatment	error_internal	0 s
get_treatment	ok	84.9 ms	get_treatment	error_internal	0 s
get_organization_treatment	ok	25.0 ms	get_treatments	error_internal	0 s
get_organization_treatment	ok	24.9 ms	get_treatments	error_internal	0 s
get_workspace_treatment	ok	21.3 ms			
get_workspace_treatment	ok	9.81 ms			
get_user_treatment	ok	5.42 ms			
get_user_treatment	ok	4.45 ms			



Migration HOWTO



Apes together strong.

Migration HOWTO

- Dedicated team:
 - 1 DevOps
 - 1 Java
 - 2 Python
 - 2 QA
- Prepare shared Helm chart and Gitlab pipelines
- Prepare list of services to migrate
- **More than 50 services in 5 months**

Migration HOWTO

- Fast and furious
- Lack of knowledge sharing

(Epic)fails



(Epic)fails

- Pay attention to resources and number of replicas
- Better to overprovision and then tune based on live metrics

Takeaways

Takeaways

- 20%-30% faster deployments
- 5% more deployments _(ツ)_/
- No more than 1 hour to start and deploy new service
- Better development experience
 - Improved CI/CD pipeline
 - Simple rollbacks
 - Easy scaling and resources management
 - Easy to deploy alerts and dashboards

Future plans

Future plans

- Dynamic environments 🔥



Thank you

101 California St #3975,
San Francisco, CA 94111,
United States